# Daniel, Eric, Tony

# Data Science Lab: Lab 7 part I (of III)

#### Submit:

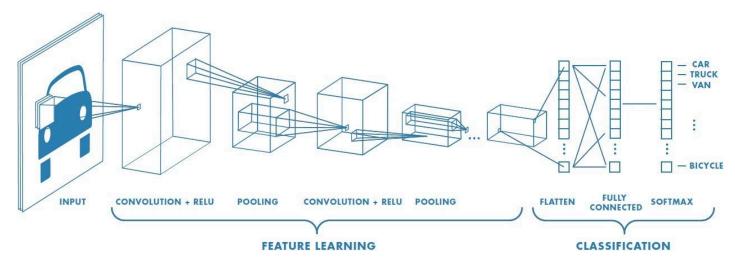
A pdf of your notebook with solutions. A link to your colab notebook or also upload your .ipynb if not working on colab.

### Goals of this Lab

There are a number of goals of this notebook:

- 1. Learning to put together a basic model beyond what we did in the previous notebook. Main emphasis: new layers, including convolution and max pooling.
- 2. Learning the basics of training.
- 3. Learning basics of loading data and visualizing.
- 4. Learning the basics of printing out a model.

Also useful to see <a href="https://pytorch.org/tutorials/beginner/blitz/cifar10\_tutorial.html">https://pytorch.org/tutorials/beginner/blitz/cifar10\_tutorial.html</a>



```
import torchvision
import torchvision.transforms as transforms
import torch
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

```
batch_size = 4
# the tutorial calls the dataloader twice -- this code defines a function
# that will do this for the train/test data.
def fetch_dataloader(batch_size, transform=None, is_train=True):
    We saw the data loaders in the previous lab.
    This creates a method for us to get (image, label) pairs.
    We can use it so that we do not have to load everything into memory
    at once.
    .....
    data = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=
    # Custom train/val split.
    indices = [i for i in range(len(data)) if (i%10 > 0) == is_train]
    data = torch.utils.data.Subset(data, indices)
    loader = torch.utils.data.DataLoader(data, batch_size=batch_size, shuffle=True, num_work
    return loader
train_transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
val_transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
data_train = fetch_dataloader(batch_size, train_transform, is_train=True)
data_val = fetch_dataloader(batch_size, val_transform, is_train=False)
→ Downloading <a href="https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz">https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz</a> to ./data/cifar-10-p
                     | 170M/170M [00:08<00:00, 20.3MB/s]
     Extracting ./data/cifar-10-python.tar.gz to ./data
     Files already downloaded and verified
```

# Problem 1 (nothing to turn in)

Read about transforms. These are routinely used when loading images. We've used a particular set of parameters. But you will see many other choices.

https://pytorch.org/vision/stable/transforms.html

### Problem 2

Figure out how to use the data loader to display the data. For example, you could look here

https://towardsdatascience.com/beginners-guide-to-loading-image-data-with-pytorch-289c60b7afec and then https://stackoverflow.com/questions/51756581/how-do-i-turn-a-pytorch-dataloader-into-a-numpy-array-to-display-image-data-with

or you could look at the pytorch CIFAR10 tutorial

https://pytorch.org/tutorials/beginner/blitz/cifar10\_tutorial.html

```
import torchvision.utils
import matplotlib.pyplot as plt
import numpy as np

data_iter = iter(data_train)
images, labels = next(data_iter)

print("Batch size:", images.size(0))  #images per match
print("Image shape:", images[0].shape)  # shape of image tensor

img = torchvision.utils.make_grid(images, nrow=4) / 2 + 0.5

npimg = img.numpy()
plt.imshow(np.transpose(npimg, (1, 2, 0)))
plt.axis('off')
plt.show()
```

Batch size: 4
Image shape: torch.Size([3, 32, 32])



# Model Implementation.

Now it is time to implement our CNN. You will recognize syntax from the previous Lab. However, whereas in the last lab we had only fully connected layers and ReLU layers, here we will use more options. In addition to fully connected and ReLU layers, we want to use:

- Convolutional layers: torch.nn.Conv2d
- Max Pooling Layers: torch.nn.MaxPool2d
- Average Pooling Layers: torch.nn.AvgPool2d

Fully connected layers primarily had 3 parameters: input size, output size, and bias.

Convolutional layers have many more parameters, as we discussed in class, In particular, recall:

- kernel\_size
- stride
- padding
- dilation

Read about these in the Pytorch documentation:

https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html

Next we see some basic code, similar in spirit to our previous colab notebook, but, critically, adding a new type of layer: a convolutional layer.

```
# We define two different models, using different syntax.
# CNNClassifier is basic code borrowed from Phil Krahenbuhl
# ConvNet comes from the pytorch tutorial linked above.
class CNNClassifier(torch.nn.Module):
    def __init__(self, input_channels, num_classes):
        Define the layer(s) needed for the model.
        Feel free to define additional input arguments.
        super().__init__()
        self.conv = nn.Conv2d(input_channels, 16, 7, 2, 3)
        self.cls = nn.Linear(16, num_classes)
        self.ReLU = nn.ReLU()
        # or could do
        # self.CNN = nn.Sequential(torch.nn.Conv2d(input_channels, 16, 7, 2, 3), torch.nn.Re
    def forward(self, x):
        Calculate the classification score (logits).
        Input:
            x (float tensor N x 3 x 32 x 32): input images
        Output:
            y (float tensor N x 10): classification scores (logits) for each class
        x = self.conv(x)
        # Add a ReLU
        x = self.ReLU(x)
        # Add global average pooling
        x = x.mean(dim=(2,3))
```

```
return self.cls(x)
    def predict(self, image):
        return self(image).argmax(1)
# From the pytorch tutorial linked above.
class ConvNet(nn.Module):
    def __init__(self):
        super(). init ()
        self.relu = nn.ReLU()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
    def forward(self, x):
        x = self.conv1(x)
        x = self.relu(x)
        x = self.pool(x)
        x = self.conv2(x)
        x = self.relu(x)
        x = self.pool(x)
        x = torch.flatten(x, 1) # flatten all dimensions except batch
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        x = self.relu(x)
        x = self.fc3(x)
        return x
# We can rewrite it equivalently this way
class ConvNet2(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = \text{torch.flatten}(x, 1) \# \text{flatten all dimensions except batch}
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

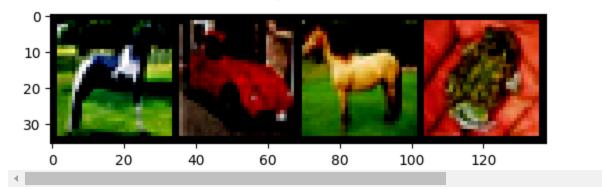
```
# This way is also equivalent
class ConvNet3(nn.Module):
    def __init__(self):
        super().__init__()
        self.relu = nn.ReLU()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
    def forward(self, x):
        x = self.pool(self.relu(self.conv1(x)))
        x = self.pool(self.relu(self.conv2(x)))
        x = \text{torch.flatten}(x, 1) \# \text{flatten all dimensions except batch}
        x = self.relu(self.fc1(x))
        x = self.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

## Model Training and Validation

```
# Now we train the first model.
# Choose a model to train
model = CNNClassifier(3, 10)
loss_function = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model.to(device)
model.train()
for epoch in range(5): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(data_train, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)
        # zero the parameter gradients
```

```
optimizer.zero_grad()
        # forward + backward + optimize
        outputs = model(inputs)
        loss = loss_function(outputs, labels)
        loss.backward()
        optimizer.step()
        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999:
                                # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
            running loss = 0.0
print('Finished Training')
\rightarrow \overline{\phantom{a}} [1, 2000] loss: 2.211
     [1, 4000] loss: 2.080
     [1, 6000] loss: 2.032
     [1, 8000] loss: 1.979
     [1, 10000] loss: 1.941
     [2, 2000] loss: 1.894
     [2, 4000] loss: 1.863
     [2, 6000] loss: 1.837
     [2, 8000] loss: 1.823
     [2, 10000] loss: 1.821
     [3, 2000] loss: 1.795
     [3, 4000] loss: 1.796
     [3, 6000] loss: 1.785
     [3, 8000] loss: 1.789
     [3, 10000] loss: 1.794
     [4, 2000] loss: 1.784
     [4, 4000] loss: 1.781
     [4, 6000] loss: 1.764
     [4, 8000] loss: 1.753
     [4, 10000] loss: 1.758
     [5, 2000] loss: 1.742
     [5, 4000] loss: 1.739
     [5, 6000] loss: 1.750
     [5, 8000] loss: 1.754
     [5, 10000] loss: 1.752
     Finished Training
# Let's see how well this trained model performs on a couple data points.
dataiter = iter(data_val)
images, labels = next(dataiter)
# print images
plt.imshow(torchvision.utils.make_grid(images).permute(1,2,0))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
# and now what our model thinks
```

₩ARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data GroundTruth: horse car horse frog
Predicted: truck ship bird dog



## Problem 3

Use your validation data loader (data\_val) to assess the accuracy on the entire data set.

If you did this correctly, you should find that your accuracy is somewhere in the 30-40% range -- not great, but a lot better than guessing, and much better than we were able to do with trees.

```
model.eval()
correct = 0
total = 0
for images, labels in data_val:
   outputs = model(images)
   val, predicted = torch.max(outputs, 1) # highest predicted class
   total += labels.size(0)
   correct += (predicted == labels).sum().item() # count number of correct
print(f'Accuracy: {100 * correct / total:.2f}%')
Accuracy: 34.86%
```

ANSWER: Accuracy: 34.86%

# Time to train the deeper model

# Now we train the second model. We see that it does quite a bit better than the first.#

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
# Choose a model to train
model2 = ConvNet() # note that this command resets the parameters
model2.to(device)
model2.train()
# Set the loss function and optimizer
loss_function = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model2.parameters(), lr=0.001, momentum=0.9)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
for epoch in range(2): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(data_train, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)
        # zero the parameter gradients
        optimizer.zero_grad()
        # forward + backward + optimize
        outputs = model2(inputs)
        loss = loss_function(outputs, labels)
        loss.backward()
        optimizer.step()
        # print statistics
        running_loss += loss.item()
                               # print every 2000 mini-batches
        if i % 2000 == 1999:
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
            running loss = 0.0
print('Finished Training')
→ [1, 2000] loss: 2.228
     [1, 4000] loss: 1.844
     [1, 6000] loss: 1.670
     [1, 8000] loss: 1.582
     [1, 10000] loss: 1.503
     [2, 2000] loss: 1.409
     [2, 4000] loss: 1.379
     [2, 6000] loss: 1.355
```

```
[2, 8000] loss: 1.332
[2, 10000] loss: 1.314
Finished Training
```

Repeat what you did for Problem 3 above: use your validation data loader (data\_val) to assess the accuracy on the entire data set.

If you did this correctly, you should find that your accuracy is somewhere in the 50-60% range. Far from perfect, but still better than how we did above.

```
model2.eval()
correct = 0
total = 0
for images, labels in data_val:
   outputs = model2(images)
   val, predicted = torch.max(outputs, 1) # highest predicted class
   total += labels.size(0)
   correct += (predicted == labels).sum().item() # count number of correct
print(f'Accuracy: {100 * correct / total:.2f}%')
Accuracy: 55.06%
```

ANSWER: Accuracy: 55.06%

### Problem 5

How many parameters does your model have? You can either compute by hand, or much better, use the summary command from torchsummary

```
from torchsummary import summary
```

Note that you'll have to figure out how to use use, and also you will have to enter the size of your input.

Also try using the command

```
for name, layer in model2.named_modules():
    print(name,layer)
```

This one is very useful when we want to download someone else's netework, and we need to know what the names of the layers are.

```
from torchsummary import summary
summary(model2, input_size=(3, 32, 32)) # found in problem 1
for name, layer in model2.named_modules():
    print(name,layer)
```

```
______
                     Output Shape
_____
       Conv2d-1
                   [-1, 6, 28, 28]
                                     456
                   [-1, 6, 28, 28]
         ReLU-2
                                      0
      MaxPool2d-3
                   [-1, 6, 14, 14]
                                      0
                  [-1, 16, 10, 10]
                                  2,416
       Conv2d-4
                  [-1, 16, 10, 10]
         ReLU-5
      MaxPool2d-6
                   [-1, 16, 5, 5]
                                      0
       Linear-7
                       [-1, 120]
         ReLU-8
                       [-1, 120]
                                  10,164
       Linear-9
                       [-1, 84]
                        [-1, 84]
        ReLU-10
       Linear-11
                        [-1, 10]
                                     850
______
Total params: 62,006
Trainable params: 62,006
```

Estimated Total Size (MB): 0.36

```
ConvNet(
  (relu): ReLU()
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in features=400, out features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)
relu ReLU()
conv1 Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
pool MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
conv2 Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
fc1 Linear(in_features=400, out_features=120, bias=True)
fc2 Linear(in_features=120, out_features=84, bias=True)
fc3 Linear(in features=84, out features=10, bias=True)
```

ANSWER: Total params = 62,006

Params size (MB): 0.24

How much can you decrease the error? Try to make your own convolutional neural network for CIFAR-10. Don't make it so big that you can't train it on Colab! You can play around with the parameters of the network and also with the parameters of training (train for longer -- more epochs, possibly using other parameters for step size, momentum, etc.).

We will give 20 lab bonus points to the team with the best accuracy.

```
# This way is also equivalent
class ConvNetCustom(nn.Module):
    def __init__(self):
        super().__init__()
        self.relu = nn.ReLU()
        self.conv1 = nn.Conv2d(3, 64, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(64, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 400)
        self.fc2_5 = nn.Linear(400, 84)
        self.fc2 = nn.Linear(400, 84)
        self.fc3 = nn.Linear(84, 10)
        self.dropout = nn.Dropout(0.5) # just cuts half the neurons
    def forward(self, x):
        x = self.pool(self.relu(self.conv1(x)))
        x = self.pool(self.relu(self.conv2(x)))
        x = \text{torch.flatten}(x, 1) \# \text{flatten all dimensions except batch}
        x = self.relu(self.fc1(x))
        x = self.dropout(x)
        x = self.relu(self.fc2(x))
        x = self.dropout(x)
        x = self.fc3(x)
        # x = self.pool(self.relu(self.conv1(x)))
        # x = self.pool(self.relu(self.conv2(x)))
        # #x = self.pool2(self.relu(self.conv3(x)))
        \# x = \text{torch.flatten}(x, 1) \# \text{flatten all dimensions except batch}
        # x = self.relu(self.fc1(x))
        # #x = self.dropout(x)
        # x = self.relu(self.fc2(x))
        # #x = self.dropout(x)
        # #x = self.relu(self.fc2_5(x))
        \# x = self.fc3(x)
        return x
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model3 = ConvNetCustom()
```

```
model3.to(device)
model3.train()
loss function = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model3.parameters(), lr=0.0015, momentum=0.85)
device = torch.device('cuda' if torch.cuda.is available() else 'cpu')
for epoch in range(7): # loop over the dataset multiple times
    running loss = 0.0
    for i, data in enumerate(data_train, 0):
       # get the inputs; data is a list of [inputs, labels]
       inputs, labels = data
       inputs = inputs.to(device)
       labels = labels.to(device)
       # zero the parameter gradients
       optimizer.zero_grad()
       # forward + backward + optimize
       outputs = model3(inputs)
       loss = loss_function(outputs, labels)
       loss.backward()
       optimizer.step()
       running loss += loss.item()
        if i % 2000 == 1999:
                                # print every 2000 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 2000))
            running loss = 0.0
print('Finished Training')
→ [1, 2000] loss: 2.183
     [1, 4000] loss: 1.904
     [1, 6000] loss: 1.742
     [1, 8000] loss: 1.676
     [1, 10000] loss: 1.611
     [2, 2000] loss: 1.542
     [2, 4000] loss: 1.487
     [2, 6000] loss: 1.451
     [2, 8000] loss: 1.424
     [2, 10000] loss: 1.397
     [3, 2000] loss: 1.341
     [3, 4000] loss: 1.336
     [3, 6000] loss: 1.325
     [3, 8000] loss: 1.278
     [3, 10000] loss: 1.304
     [4, 2000] loss: 1.226
```

[4, 4000] loss: 1.228

```
[4, 6000] loss: 1.222
     [4, 8000] loss: 1.212
     [4, 10000] loss: 1.196
     [5, 2000] loss: 1.166
     [5, 4000] loss: 1.158
     [5, 6000] loss: 1.161
     [5, 8000] loss: 1.152
     [5, 10000] loss: 1.147
     [6, 2000] loss: 1.088
     [6, 4000] loss: 1.102
     [6, 6000] loss: 1.115
     [6, 8000] loss: 1.105
     [6, 10000] loss: 1.095
     [7, 2000] loss: 1.031
     [7, 4000] loss: 1.045
     [7, 6000] loss: 1.068
     [7, 8000] loss: 1.070
     [7, 10000] loss: 1.077
     Finished Training
model3.eval()
correct = 0
total = 0
for images, labels in data_val:
  outputs = model3(images)
  val, predicted = torch.max(outputs, 1) # highest predicted class
  total += labels.size(0)
  correct += (predicted == labels).sum().item() # count number of correct
print(f'Accuracy: {100 * correct / total:.2f}%')
→ Accuracy: 65.82%
```

ANSWER: Accuracy: 65.82%

#### Data Science Lab: Lab 7 part II (of III)

Submit:

A pdf of your notebook with solutions. A link to your colab notebook or also upload your .ipynb if not working on colab.

#### Goals of this Lab: Transfer Learning

#### https://pytorch.org/tutorials/beginner/transfer\_learning\_tutorial.html

In the previous notebook, we learned how to put together various different types of layers to build a convolutional neural network (CNN) to classify CIFAR-10. Now we are going to learn about the principle of transfer learning.

We will create 2 data sets: one with the first five labels of CIFAR-10 (plane, car, bird, cat, deer), and one with the remaining 5. We will see how training on the second data set, can actually help us with the first.

Note: because CIFAR-10 is a relatively small data set, the numbers are not overly compelling. But we can greatly amplify this if we use larger data sets.

#### Problem 1.

Make train and validation (testing) data loaders for the first five labels, and the second five labels.

You should have 4 dataloaders when done.

```
# First we download the data. This is quite similar to what we've seen before.
# The main difference will be that we need to make two different loaders,
# one for the first five labels, and one for the remaining five.
import torchvision
import torchvision.transforms as transforms
import torch
import torch.nn as nn
import torch.nn.functional as F
torch.manual_seed(42)
classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
batch_size = 4
# the tutorial calls the dataloader twice -- this code defines a function
# that will do this for the train/test data.
def fetch_dataloader(batch_size, label_subset,transform=None, is_train=True):
    Loads data from disk and returns a data_loader.
    A DataLoader is similar to a list of (image, label) tuples.
    You do not need to fully understand this code to do this assignment, we're happy to explain though.
    data = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
    # replaced TO DOs with following 5 lines
    indices = []
    for i, sample in enumerate(data):
        if sample[1] in label_subset:
            indices.append(i)
    data = torch.utils.data.Subset(data, indices)
    loader = torch.utils.data.DataLoader(data, batch_size=batch_size, shuffle=True, num_workers=2)
    return loader
train transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms. Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])\\
val_transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```

```
TO DO
Complete these four using the function you wrote above, to get the 4 data loaders you need:
two for labels 1-5, and two for labels 6-10.
data15_train = fetch_dataloader(batch_size, [0, 1, 2, 3, 4], transform=train_transform, is_train=True)
data15_val = fetch_dataloader(batch_size, [0, 1, 2, 3, 4], transform=val_transform, is_train=False)
data610_train = fetch_dataloader(batch_size, [5, 6, 7, 8, 9], transform=train_transform, is_train=True)
data610_val = fetch_dataloader(batch_size, [5, 6, 7, 8, 9], transform=val_transform, is_train=False)
Files already downloaded and verified
     Files already downloaded and verified
     Files already downloaded and verified
     Files already downloaded and verified
# Now let's visualize a few images, to see if this worked.
# Same exact syntax as in the previous notebook.
# Nothing to do here, except run it to make sure it worked.
import matplotlib.pyplot as plt
import numpy as np
from torchvision.utils import make_grid
%matplotlib inline
def imshow(img):
   img = img / 2 + 0.5
                            # unnormalize
   npimg = img.numpy()
   plt.imshow(np.transpose(npimg, (1, 2, 0)))
   plt.show()
# get some random training images
dataiter = iter(data610_train)
images, labels = next(dataiter)
# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join('%5s' % classes[labels[j]] for j in range(batch_size)))
→
      10
      20
      30
                  20
                            40
                                     60
                                              80
                                                       100
                                                                 120
    chin truck chin horca
```

```
# that we got from the Pytorch ConvNet tutorial.
class ConvNet(nn.Module):
   def __init__(self):
       super().__init__()
        self.relu = nn.ReLU()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
   def forward(self, x):
       x = self.conv1(x)
        x = self.relu(x)
       x = self.pool(x)
       x = self.conv2(x)
       x = self.relu(x)
       x = self.pool(x)
       x = \text{torch.flatten}(x, 1) \# \text{flatten all dimensions except batch}
       x = self.fc1(x)
        x = self.relu(x)
```

# We define the same model from the last notebook

```
x = self.fc2(x)
        x = self.relu(x)
        x = self.fc3(x)
        return x
# We train the model on the first five labels to see how well we do after a fixed amount of training.
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
# Choose a model to train
model1 = ConvNet() # note that this command resets the parameters
model1.to(device)
model1.train()
# Set the loss function and optimizer
loss_function = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model1.parameters(), lr=0.001, momentum=0.9)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
for epoch in range(1): # a single pas through the data. Don't change this.
    running_loss = 0.0
    for i, data in enumerate(data15_train, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)
        # zero the parameter gradients
        optimizer.zero_grad()
        # forward + backward + optimize
        outputs = model1(inputs)
        loss = loss_function(outputs, labels)
        loss.backward()
        optimizer.step()
        # print statistics
        running_loss += loss.item()
        if i % 200 == 199:  # print every 200 mini-batches
            print('[%d, %5d] loss: %.3f' %
                   (epoch + 1, i + 1, running_loss / 200))
             running_loss = 0.0
print('Finished Training')
           200] loss: 2.020
 <del>____</del> [1,
     [1, 400] loss: 1.613
          600] loss: 1.574
     [1,
     [1, 800] loss: 1.484
     [1, 1000] loss: 1.406
     [1, 1200] loss: 1.361
[1, 1400] loss: 1.289
     [1, 1600] loss: 1.274
[1, 1800] loss: 1.286
     [1, 2000] loss: 1.254
     [1, 2200] loss: 1.198
     [1, 2400] loss: 1.164
     [1, 2600] loss: 1.196
     [1, 2800] loss: 1.138
[1, 3000] loss: 1.166
     [1, 3200] loss: 1.159
     [1, 3400] loss: 1.081
[1, 3600] loss: 1.104
     [1, 3800] loss: 1.105
     [1, 4000] loss: 1.084
[1, 4200] loss: 1.097
     [1, 4400] loss: 1.046
[1, 4600] loss: 1.039
     [1, 4800] loss: 1.042
     [1, 5000] loss: 1.049
     [1, 5200] loss: 1.046
     [1, 5400] loss: 1.010
     [1, 5600] loss: 0.982
     [1, 5800] loss: 1.056
```

```
[1, 6000] loss: 1.033
[1, 6200] loss: 1.008
Finished Training
```

In the cell above, you trained your model on labels 1-5 for one single epoch.

How well does this do on the validation data set?

This is nearly identical to problems you solved in part I.

```
ANSWER: Model 1 accuracy: 61.94%
```

```
model1.eval()
correct = 0
total = 0
for images, labels in data15_val:
 outputs = model1(images)
 val, predicted = torch.max(outputs, 1) # highest predicted class
 total += labels.size(0)
 correct += (predicted == labels).sum().item() # count number of correct
print(f'Accuracy: {100 * correct / total:.2f}%')
→ Accuracy: 61.94%
# Now we train the model on the last 5 labels to use for transfer learning.
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
# Choose a model to train
model2 = ConvNet() # note that this command resets the parameters
model2.to(device)
model2.train()
# Set the loss function and optimizer
loss_function = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model2.parameters(), lr=0.001, momentum=0.9)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
for epoch in range(6): # loop over the dataset multiple times
    running_loss = 0.0
    for i, data in enumerate(data610_train, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        labels = labels-5 # hack b/c labels are 2-9 instead of 0-7
        inputs = inputs.to(device)
        labels = labels.to(device)
        # zero the parameter gradients
        optimizer.zero_grad()
        # forward + backward + optimize
        outputs = model2(inputs)
        loss = loss_function(outputs, labels)
        loss.backward()
        optimizer.step()
        # print statistics
        running_loss += loss.item()
        if i % 500 == 499:
                            # print every 500 mini-batches
            print('[%d, %5d] loss: %.3f' %
                  (epoch + 1, i + 1, running_loss / 500))
            running_loss = 0.0
print('Finished Training')
<del>_</del>
```

```
[2, 4000] 1055: 0./3/
[2, 4500] loss: 0.722
    5000] loss: 0.717
[2, 5500] loss: 0.681
[2, 6000] loss: 0.678
[3,
     500] loss: 0.605
[3, 1000] loss: 0.654
[3,
    1500] loss: 0.644
    20001 loss: 0.616
Γ3,
[3,
    2500] loss: 0.620
[3,
    3000] loss: 0.597
    35001 loss: 0.637
[3,
[3,
    4000] loss: 0.604
    4500]
          loss: 0.624
    5000] loss: 0.595
[3,
[3,
    5500] loss: 0.607
[3,
    6000] loss: 0.616
     500] loss: 0.522
Γ4,
    10001 loss: 0.540
۲4,
[4,
    1500] loss: 0.532
[4,
    2000] loss: 0.558
    2500] loss: 0.554
Γ4,
Γ4,
    30001 loss: 0.541
[4,
    3500] loss: 0.528
[4,
    4000] loss: 0.559
[4,
    45001 loss: 0.542
    5000] loss: 0.593
     5500]
          loss: 0.561
[4,
[4,
    6000] loss: 0.537
[5,
     500] loss: 0.492
    1000] loss: 0.501
    1500] loss: 0.481
[5,
    20001 loss: 0.476
[5,
[5,
    2500] loss: 0.514
[5,
    3000] loss: 0.466
[5,
    35001 loss: 0.490
    4000] loss: 0.500
[5,
    4500] loss: 0.501
[5,
    5000] loss: 0.479
[5,
    55001 loss: 0.492
[5,
[5,
    6000] loss: 0.477
[6,
     500]
          loss: 0.429
[6, 1000] loss: 0.459
[6,
    1500] loss: 0.401
[6,
    2000] loss: 0.449
    2500] loss: 0.426
ſ6,
[6,
    3000] loss: 0.417
[6,
    35001 loss: 0.466
    4000] loss: 0.474
[6,
    45001 loss: 0.464
[6,
[6,
    5000] loss: 0.485
    5500] loss: 0.451
    6000] loss: 0.424
[6,
Finished Training
```

Now we see how we can use this model for the first 5 labels.

```
# We define a new model to fine-tune
model3 = model2
```

We have a new model: model3. This model has the same weights as model2, hence it is good at classifying labels, 6-10, and does not know about labels 1-5. We will have to reset and retrain the last layer.

#### Problem 3

Let's pretend that didn't build model3 ourselves. In this case, to reset the last layer, we would need to figure out what its name is (we know it's fc3, but we're pretending we don't).

Find a command that will tell you the name and type of the last layer. Print them out.

```
for name, layer in model3.named_modules():
    print(name, layer)

ConvNet(
          (relu): ReLU()
          (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
          (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
```

```
(conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)
relu ReLU()
conv1 Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
pool MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
conv2 Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
fc1 Linear(in_features=400, out_features=120, bias=True)
fc2 Linear(in_features=120, out_features=84, bias=True)
fc3 Linear(in_features=84, out_features=10, bias=True)
```

ANSWER: We can see above that the last layer is called fc3 and is a linear layer

#### Problem 4

If we did the above correctly, then we learned that the last layer is called fc3 and it is a linear layer.

We want to redefine it giving a command like

```
model3.fc3 = nn.Linear(input_features,output_features)
```

- What should be the value of output\_features for our new last layer that we are defining? The answer to this comes from the problem we
  want to solve.
- What is the number of input features that go into this last layer? You'll have to look at the size of the layer defined in the network. Again, pretend we didn't make the network, and we just downloaded it.

So in other words, I'm asking you to use commands to find out how many input features there are for the last layer which (we know from above) is called fc3.

ANSWER: output\_features should be 5 because we are predicting for 5 classes. The number of input features is the number of output features of the previous layer fc2.

#### Problem 5

Now redefine the last layer. You will do this with a command like

```
model3.fc3 = # TO DO
```

Note how important it is for us to know what the names of the layers are.

```
model3.fc3 = nn.Linear(model3.fc2.out features, 5)
model3 = model3.to(device)
model3.train()
# Let's again check the layers and see the change we made
# (doesn't actually look different since we split 5 and 5)
for name, layer in model2.named_modules():
 print(name,layer)
     ConvNet(
       (relu): ReLU()
       (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
       (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
       (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
       (fc1): Linear(in_features=400, out_features=120, bias=True)
       (fc2): Linear(in_features=120, out_features=84, bias=True)
       (fc3): Linear(in_features=84, out_features=5, bias=True)
     conv1 Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
     pool MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
     conv2 Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
     fc1 Linear(in_features=400, out_features=120, bias=True)
     fc2 Linear(in_features=120, out_features=84, bias=True)
     fc3 Linear(in_features=84, out_features=5, bias=True)
```

```
# Set the optimizer (loss function already set)
optimizer = torch.optim.SGD(model3.parameters(), lr=0.001, momentum=0.9)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

Check that model3 is no good for classes 1-5. (It shouldn't be any good! We just added a randomly initialized last layer). The accuracy should be close to 20%, which is how good random guessing would be for 5 labels.

Evaluate its accuracy on the validation set for labels 1-5.

ANSWER: Model 3 accuracy before training: 16.51%

```
model3.eval()
correct = 0
total = 0
for images, labels in data15_val:
 outputs = model3(images)
 val, predicted = torch.max(outputs, 1) # highest predicted class
 total += labels.size(0)
 correct += (predicted == labels).sum().item() # count number of correct
print(f'Accuracy: {100 * correct / total:.2f}%')
→ Accuracy: 16.51%
\mbox{\#}\mbox{Now} we train model3 on the first five labels, and evaluate accuracy.
# Nothing for you to do here but run this cell.
# Note that we are also just using one single pass, to make it a fair
# comparison, and so that we can try to see the benefit of transfer learning.
for epoch in range(1): # one single pass -- don't change this
    running_loss = 0.0
    for i, data in enumerate(data15_train, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.to(device)
        labels = labels.to(device)
        # zero the parameter gradients
        optimizer.zero_grad()
        # forward + backward + optimize
        outputs = model3(inputs)
        loss = loss_function(outputs, labels)
        loss.backward()
        optimizer.step()
        # print statistics
        running_loss += loss.item()
        if i % 200 == 199:
                            # print every 200 mini-batches
            print('[%d, %5d] loss: %.3f' \%
                  (epoch + 1, i + 1, running_loss / 200))
            running_loss = 0.0
print('Finished Training')
<del>____</del> [1,
           200] loss: 1.182
           400] loss: 1.043
     [1,
           600] loss: 1.036
     [1,
     [1, 800] loss: 1.02/
[1, 1000] loss: 1.002
     [1, 1200] loss: 0.991
     [1, 1400] loss: 0.920
     [1, 1600] loss: 0.917
     [1, 1800] loss: 0.950
     [1, 2000] loss: 0.935
     [1, 2200] loss: 0.924
     [1, 2400] loss: 0.902
     [1, 2600] loss: 0.896
[1, 2800] loss: 0.862
     [1, 3000] loss: 0.888
     [1, 3200] loss: 0.808
```

```
[1, 3400] loss: 0.869
[1, 3600] loss: 0.915
[1, 3800] loss: 0.925
[1, 4000] loss: 0.846
[1, 4200] loss: 0.846
[1, 4600] loss: 0.862
[1, 4800] loss: 0.799
[1, 5000] loss: 0.829
[1, 5200] loss: 0.839
[1, 5600] loss: 0.839
[1, 5600] loss: 0.890
[1, 6000] loss: 0.753
[1, 6200] loss: 0.815
Finished Training
```

Evaluate the accuracy of model3 (after 1 epoch of training) on the validation set. Compare with the accuracy you found above in Problem 2.

ANSWER: Model 3 accuracy after training: 70.07% which is about 8% higher than the one found in problem 2

```
model3.eval()
correct = 0
total = 0
for images, labels in data15_val:
    outputs = model3(images)
    val, predicted = torch.max(outputs, 1) # highest predicted class
    total += labels.size(0)
    correct += (predicted == labels).sum().item() # count number of correct
print(f'Accuracy: {100 * correct / total:.2f}%')
Accuracy: 70.07%
```

# Data Science Lab: Lab 7 part III (of III)

Submit:

A pdf of your notebook with solutions. A link to your colab notebook or also upload your .ipynb if not working on colab.

## Goals of this Lab: Fine Tuning and Transfer Learning

This assignment completes our segment on convolutional networks, computer vision and transfer learning. The goal of this colab notebook and lab is to get down the basics of Transfer learning, and harnessing the power of pre-trained notebooks, in a more substantial setting than CIFAR-10. Along the way, we get experience with creating and labeling data sets, and a number of other Python tools.

Specifically, we do the following:

- In this colab notebook, we create our own labeled image data set, with calls to the Bing API. We use this to download images with user-specified labels, into train and test directories.
- · We preprocess all the images and use the directory names as lables to create our training and testing data sets.
- Then we download a pre-trained convolutional neural network from Pytorch. There is wide selection here. These have been trained on ImageNet. See references below.
- Then we add a new last layer, and train. Note that the last layer we add has to have the right size, namely, the number of classes in our data.
- Things to play with: Choosing different pre-trained models; Fine tuning the entire network vs freezing and fine
  tuning only the last layer or layers; Possibly adding more of our own layers at the end. (See the last optional
  problem at the end).

#### Some references:

The bing image downloader package from here <a href="https://pypi.org/project/bing-image-downloader/">https://pypi.org/project/bing-image-downloader/</a> <a href="https://pypi.org/project/bing-image-downloader/">https://pypi.org/project/bing-image-downloader/</a>)

# The following will allow us to call the Bing Image Downloader

Requirement already satisfied: bing-image-downloader in c:\users\tonys\anaconda3\envs\gpu\_pytorchenv\lib\site-packages (1.1.2)

```
In [6]: ▶
            1 from __future__ import print_function, division
             3 import torch
             4 import torch.nn as nn
             5 import torch.optim as optim
             6 from torch.optim import lr_scheduler
             7 import numpy as np
             8 import torchvision
             9 from torchvision import datasets, models, transforms
            10 import time
            11 import os
            12 import copy
            13 import requests
            14 from bing_image_downloader import downloader
            15 import matplotlib.pyplot as plt
            16 import matplotlib.image as mpimg
            17 import random
            18 import math
            19
            20 plt.ion() # interactive mode
```

Out[6]: <contextlib.ExitStack at 0x1fb263f53a0>

# Downloading the images we need

```
In [4]:
              1
                # These are the functions we need for asking for query terms, number of images
                # and the fraction of train/test split
              3
                def get_query_terms():
                     """Ask the user for a list of query terms."""
              4
              5
                     queries = input("Enter query terms separated by commas: ").split(',')
              6
                     return [query.strip() for query in queries]
              7
              8
              9
                def get_positive_integers():
             10
                     """Ask the user for a list of positive integers."""
             11
                     while True:
                         numbers = input("Enter positive integers corresponding to each query, separated
             12
             13
                        try:
             14
                             # Convert string inputs to integers
             15
                             numbers = [int(num.strip()) for num in numbers]
             16
             17
                             # Check if all numbers are positive
             18
                             if all(num > 0 for num in numbers):
             19
                                 return numbers
             20
                             else:
                                 print("All numbers must be positive. Please try again.")
             21
             22
                         except ValueError:
                             print("Invalid input. Please enter integers only.")
             23
             24
             25
             26 def get_train_test_proportion():
             27
                     """Ask the user for a number between 0 and 1 and ensure it is in that range."""
             28
                     while True:
             29
                         try:
                             # Ask the user for a number and try to convert it to a float
             30
             31
                             num = float(input("What fraction of the data should be training? For exampl
             32
             33
                             # Check if the number is between 0 and 1 (inclusive)
             34
                             if 0 <= num <= 1:
             35
                                 return num
             36
                             else:
             37
                                 print("The number must be between 0 and 1. Please try again.")
             38
                         except ValueError:
                             print("Invalid input. Please enter a valid number between 0 and 1.")
             39
```

Run the above code, and the code below, to make your own data set. Choose terms that aren't already one of the 1000 classes in the Imagenet dataset.

I recommend starting small: 2 or 3 different terms, and maybe 100-200 images max from each.

```
In [8]:
             1 import shutil
             3 queries = get_query_terms()
             4 # And the number of each
             5 numbers = get positive integers()
             6 # And the fraction of training data
             7 frac_train = get_train_test_proportion()
             8 # Check if number of queries matches number of positive integers
             9 if len(queries) != len(numbers):
             10
                  print("Number of queries and integers do not match!")
             11
             12 for query, n in zip(queries, numbers):
             13
                    # Download all images to a common directory
             14
                    downloader.download(query, limit=n, output_dir='dataset/all_images', adult_filter_o
             15
             16
                    # Directory paths
             17
                    source_dir = f'dataset/all_images/{query}'
                    train_dir = f'dataset/train/{query}'
             18
             19
                    test_dir = f'dataset/test/{query}'
             20
                    # Create train and test directories if they don't exist
             21
             22
                    os.makedirs(train_dir, exist_ok=True)
             23
                    os.makedirs(test_dir, exist_ok=True)
             24
             25
                    # Get all downloaded files
             26
                    all_files = [os.path.join(source_dir, f) for f in os.listdir(source_dir)]
             27
                    random.shuffle(all_files) # Shuffle to randomize distribution
             28
             29
                    # Determine split counts
             30
                    n_train = int(math.floor(len(all_files) * frac_train))
             31
                    train_files = all_files[:n_train]
             32
                    test_files = all_files[n_train:]
             33
             34
                    # Move files to respective directories
             35
                    for f in train_files:
             36
                        shutil.move(f, train_dir)
             37
                    for f in test_files:
             38
                         shutil.move(f, test_dir)
             39
             40
                    # Optionally, remove the source directory if it's now empty
             41
                    if not os.listdir(source_dir):
             42
                        os.rmdir(source_dir)
             43
```

Enter query terms separated by commas: 'vintage\_typewriter', 'antique\_compass'
Enter positive integers corresponding to each query, separated by commas: 150,150
What fraction of the data should be training? For example, enter 0.75 for 75 percent: 0.8
[%] Downloading Images to C:\Users\tonys\Desktop\Fall 2024\Data Science Lab\Lab homework\d ataset\all\_images\'vintage\_typewriter'

[Error]Invalid image, not saving https://preview.redd.it/up2w6modfti51.jpg?auto=webp&s=7e393f006607bf946bd5f741b9f8819de5a30070 (https://preview.redd.it/up2w6modfti51.jpg?auto=webp&s=7e393f006607bf946bd5f741b9f8819de5a30070)

- [!] Issue getting: https://preview.redd.it/up2w6modfti51.jpg?auto=webp&s=7e393f006607bf946bd5f741b9f8819de5a30070 (https://preview.redd.it/up2w6modfti51.jpg?auto=webp&s=7e393f006607bf946bd5f741b9f8819de5a30070)
- [!] Error:: Invalid image, not saving https://preview.redd.it/up2w6modfti51.jpg?auto=webp& amp;s=7e393f006607bf946bd5f741b9f8819de5a30070 (https://preview.redd.it/up2w6modfti51.jpg?auto=webp&s=7e393f006607bf946bd5f741b9f8819de5a30070)

[Error]Invalid image, not saving https://i.redd.it/2zl58ghgfl421.jpg (https://i.redd.it/2zl58ghgfl421.jpg)

- [!] Issue getting: https://i.redd.it/2zl58ghgfl421.jpg (https://i.redd.it/2zl58ghgfl421.jp
  g)
- [!] Error:: Invalid image, not saving https://i.redd.it/2z158ghgfl421.jpg (https://i.redd. it/2z158ghgfl421.jpg)
- [!] Issue getting: https://c1.wallpaperflare.com/preview/962/977/65/typewriter-vintage-old -vintage-typewriter.jpg (https://c1.wallpaperflare.com/preview/962/977/65/typewriter-vintage-old-vintage-typewriter.jpg)
- [!] Error:: <urlopen error [SSL] record layer failure (\_ssl.c:1133)>
  [Error]Invalid image, not saving https://images.squarespace-cdn.com/content/v1/5ca223529d4
  14920186973a5/1678120691582-VWX48PM0BMGJCUV06BXA/Remington+Model+5+Vintage+Manual+Typewrit
  er.jpg (https://images.squarespace-cdn.com/content/v1/5ca223529d414920186973a5/16781206915
  82-VWX48PM0BMGJCUV06BXA/Remington+Model+5+Vintage+Manual+Typewriter.jpg)
- [!] Issue getting: https://images.squarespace-cdn.com/content/v1/5ca223529d414920186973a5/1678120691582-VWX48PM0BMGJCUV06BXA/Remington+Model+5+Vintage+Manual+Typewriter.jpg (https://images.squarespace-cdn.com/content/v1/5ca223529d414920186973a5/1678120691582-VWX48PM0BMGJCUV06BXA/Remington+Model+5+Vintage+Manual+Typewriter.jpg)
- [!] Error:: Invalid image, not saving https://images.squarespace-cdn.com/content/v1/5ca223 529d414920186973a5/1678120691582-VWX48PM0BMGJCUV06BXA/Remington+Model+5+Vintage+Manual+Typewriter.jpg (https://images.squarespace-cdn.com/content/v1/5ca223529d414920186973a5/167812 0691582-VWX48PM0BMGJCUV06BXA/Remington+Model+5+Vintage+Manual+Typewriter.jpg)

[Error]Invalid image, not saving https://preview.redd.it/ypnikseb7u431.jpg?auto=webp&s =0c7b3fcd9087ed8cb6c9eb257e68e9d8f84b031e (https://preview.redd.it/ypnikseb7u431.jpg?auto=webp&s=0c7b3fcd9087ed8cb6c9eb257e68e9d8f84b031e)

- [!] Issue getting: https://preview.redd.it/ypnikseb7u431.jpg?auto=webp&s=0c7b3fcd9087ed8cb6c9eb257e68e9d8f84b031e (https://preview.redd.it/ypnikseb7u431.jpg?auto=webp&s=0c7b3fcd9087ed8cb6c9eb257e68e9d8f84b031e)
- [!] Error:: Invalid image, not saving https://preview.redd.it/ypnikseb7u431.jpg?auto=webp&s=0c7b3fcd9087ed8cb6c9eb257e68e9d8f84b031e (https://preview.redd.it/ypnikseb7u431.jpg?auto=webp&s=0c7b3fcd9087ed8cb6c9eb257e68e9d8f84b031e)
- [%] Done. Downloaded 150 images.
- [%] Downloading Images to C:\Users\tonys\Desktop\Fall 2024\Data Science Lab\Lab homework\d ataset\all images\'antique compass'
- [!] Issue getting: http://www.geodus.com/globe-map/tresgrand/am\_CO012.jpg (http://www.geodus.com/globe-map/tresgrand/am\_CO012.jpg)
- [!] Error:: <urlopen error [SSL: CERTIFICATE\_VERIFY\_FAILED] certificate verify failed: una ble to get local issuer certificate ( ssl.c:1133)>
- [!] Issue getting: https://churchantiques.com/wp-content/uploads/2021/01/IMG\_20201123\_1636 38.jpg (https://churchantiques.com/wp-content/uploads/2021/01/IMG\_20201123\_163638.jpg)
- [!] Error:: HTTP Error 404: Not Found
- [!] Issue getting: https://www.craple.com/wp-content/uploads/2023/03/Collecting-Antique-Co

```
mpasses.webp (https://www.craple.com/wp-content/uploads/2023/03/Collecting-Antique-Compass
es.webp)
[!] Error:: Remote end closed connection without response
[!] Issue getting: http://www.scientificcollectables.com/PHOTOS (http://www.scientificcoll
ectables.com/PHOTOS) 2019/sc2302c.jpg
[!] Error:: URL can't contain control characters. '/PHOTOS 2019/sc2302c.jpg' (found at lea
st ' ')
[!] Issue getting: https://media.istockphoto.com/photos/antique-compass-picture-id17354383
7 (https://media.istockphoto.com/photos/antique-compass-picture-id173543837)
[!] Error:: HTTP Error 400: Bad Request
[!] Issue getting: https://www.crushpixel.com/big-static11/preview4/old-vintage-compass-on
-ancient-648094.jpg (https://www.crushpixel.com/big-static11/preview4/old-vintage-compass-
on-ancient-648094.jpg)
[!] Error:: HTTP Error 403: Forbidden
[Error]Invalid image, not saving http://www.windandweather.com/medias/sys master/images/im
ages/hb7/hc0/8838039896094/IN7303SInglesx.jpg (http://www.windandweather.com/medias/sys_ma
ster/images/images/hb7/hc0/8838039896094/IN7303SInglesx.jpg)
[!] Issue getting: http://www.windandweather.com/medias/sys master/images/images/hb7/hc0/8
838039896094/IN7303SInglesx.jpg (http://www.windandweather.com/medias/sys_master/images/im
ages/hb7/hc0/8838039896094/IN7303SInglesx.jpg)
[!] Error:: Invalid image, not saving http://www.windandweather.com/medias/sys master/imag
es/images/hb7/hc0/8838039896094/IN7303SInglesx.jpg (http://www.windandweather.com/medias/s
ys master/images/images/hb7/hc0/8838039896094/IN7303SInglesx.jpg)
[Error]Invalid image, not saving https://www.nauticalia-trade-sales.com/index.php?route=pr
oduct/product/saveImage&image=image/catalog/products_images/4429.jpg (https://www.naut
icalia-trade-sales.com/index.php?route=product/product/saveImage&image=image/catalog/p
roducts_images/4429.jpg)
[!] Issue getting: https://www.nauticalia-trade-sales.com/index.php?route=product/product/
saveImage&image=image/catalog/products_images/4429.jpg (https://www.nauticalia-trade-s
ales.com/index.php?route=product/product/saveImage&image=image/catalog/products image
s/4429.ipg)
[!] Error:: Invalid image, not saving https://www.nauticalia-trade-sales.com/index.php?rou
te=product/product/saveImage&image=image/catalog/products images/4429.jpg (https://ww
w.nauticalia-trade-sales.com/index.php?route=product/product/saveImage&image=image/cat
alog/products_images/4429.jpg)
[!] Issue getting: https://antiqueitemstore.com/wp-content/uploads/2020/04/5-6-644x1024.jp
g (https://antiqueitemstore.com/wp-content/uploads/2020/04/5-6-644x1024.jpg)
[!] Error:: Remote end closed connection without response
```

# Let's make sure we downloaded enough of each

[%] Done. Downloaded 150 images.

The code above is just making calls to the Bing API, and as you'll see there are a number of errors. So we should make sure that we have in fact downloaded the images we need. Print the number of files

```
In [10]:
              1
                 import os
              3
                 def count_files_in_directory(directory, prefix=""):
              4
              5
                     Recursively counts the number of files in each directory and subdirectory.
              6
              7
              8
                     - directory: The directory path to start counting from.
              9
                     - prefix: A string used for indentation to visualize the folder structure.
              10
              11
                     num_files = sum([len(files) for r, d, files in os.walk(directory)])
              12
                     print(f"{prefix}{os.path.basename(directory)}: {num_files} files")
              13
              14
                     for subdir in next(os.walk(directory))[1]: # List subdirectories of the current di
              15
                         path = os.path.join(directory, subdir) # Full path of the subdirectory
                         count_files_in_directory(path, prefix + " ") # Recursively count in this subd
              16
              17
              18 # Path to the dataset directory
              19 dataset_dir = 'C:\\Users\\tonys\\Desktop\\Fall 2024\\Data Science Lab\\Lab homework\\da
              20
              21 # Start the recursive count
              22 count_files_in_directory(dataset_dir)
             dataset: 312 files
               all images: 0 files
               test: 63 files
                 'antique_compass': 30 files
                 'vintage_typewriter': 33 files
```

# Preprocessing the data

'antique\_compass': 120 files
'vintage\_typewriter': 129 files

train: 249 files

Now that we've downloaded the data we want, we do some basic preprocessing with functions from Torchvision.

For the training data, we normalize and also use some data augmentation.

For the validation set, we just normalize.

# **Problem 1 (Nothing to turn in)**

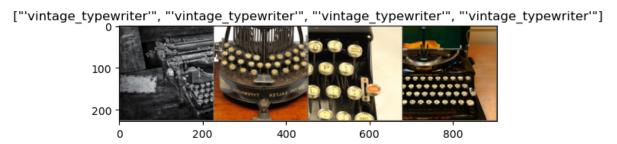
Read about this here: <a href="https://pytorch.org/vision/stable/transforms.html">https://pytorch.org/vision/stable/transforms.html</a> (<a href="https://pytorch.org/vision/stable/transforms.html">https://pytorch.org/vision/stable/transforms.html</a>)

# Make dataloaders out of our training and testing data

```
In [11]:
              1 # Data augmentation and normalization for training
              2 # Just normalization for validation
              4 # The directory that contains the data
              5 data_dir = 'dataset'
              7 # Now we apply the usual transformations
                 data_transforms = {
              9
                     'train': transforms.Compose([
              10
                         transforms.RandomResizedCrop(224),
              11
                         transforms.RandomHorizontalFlip(),
                         transforms.ToTensor(),
              12
             13
                         transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
             14
                     ]),
             15
                      'test': transforms.Compose([
             16
                         transforms.Resize(256),
             17
                         transforms.CenterCrop(224),
             18
                         transforms.ToTensor(),
                         transforms.Normalize([0.485, 0.456, 0.406], [0.229, 0.224, 0.225])
             19
              20
                     ]),
              21 }
              22
                 image_datasets = {x: datasets.ImageFolder(os.path.join(data_dir, x),
              23
                                                           data_transforms[x])
                                   for x in ['train', 'test']}
              24
              25 dataloaders = {x: torch.utils.data.DataLoader(image_datasets[x], batch_size=4,
              26
                                                              shuffle=True, num_workers=2)
              27
                               for x in ['train', 'test']}
              28 dataset_sizes = {x: len(image_datasets[x]) for x in ['train', 'test']}
              29 class_names = image_datasets['train'].classes
              30
              31 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

Write code that calls the dataloaders and displays some number of images, with their true labels.

```
In [12]:
               1
                  import matplotlib.pyplot as plt
               2
               3
                  def imshow(inp, title=None):
                      """Imshow for Tensor."""
               4
               5
                      inp = inp.numpy().transpose((1, 2, 0))
               6
                      mean = np.array([0.485, 0.456, 0.406])
               7
                      std = np.array([0.229, 0.224, 0.225])
               8
                      inp = std * inp + mean
               9
                      inp = np.clip(inp, 0, 1)
              10
                      plt.imshow(inp)
              11
                      if title is not None:
              12
                          plt.title(title)
              13
                      plt.pause(0.001) # pause a bit to update plots
              14
                  # Get a batch of training data
              15
              16 | inputs, classes = next(iter(dataloaders['train']))
              17
              18 #Make a grid from batch
              19
                  out = torchvision.utils.make_grid(inputs)
              20
                 # Visualize
              21
                  imshow(out, title=[class_names[x] for x in classes])
```



# Download a pre-trained model

We will now download a pretrained model that has been trained on Imagenet. I recomend starting with Resnet18. See here for other larger or smaller pre-trained modfels <a href="https://pytorch.org/vision/stable/models.html">https://pytorch.org/vision/stable/models.html</a> (<a href="https://pytorch.org/vision/stable/models.html">https://pytorch.org/vision/stable/models.html</a>)

```
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to C:\Users\tony s/.cache\torch\hub\checkpoints\resnet18-f37072fd.pth

100%| 44.7M [00:01<00:00, 40.2MB/s]
```

### **Problem 3**

Redefine the last layer of your neural network (called model\_ft) to be a linear (fully connected layer) whose input size is whatever the input size is of the the current last layer, and whose output size is the number of labels your new data set has.

You need to figure out how to do this. It's not too difficult. You need will use similar commands as you used in the last 2 homeworks in order to define a linear layer. The key will be to find out what the last layer is called, and also to find

# **Training**

- We are going to fine-tune by training all the layers.
- We can also freeze the old layers, and only update the last layers that we added. (This is not implemented below.)

```
In [28]: ▶
               1 | def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
               2
                      since = time.time()
               3
               4
                      best_model_wts = copy.deepcopy(model.state_dict())
               5
                      best_acc = 0.0
               6
               7
                      for epoch in range(num_epochs):
               8
                          print('Epoch {}/{}'.format(epoch, num_epochs - 1))
               9
                          print('-' * 10)
              10
              11
                          # Each epoch has a training and validation phase
              12
                          for phase in ['train', 'test']:
                              if phase == 'train':
              13
              14
                                  model.train() # Set model to training mode
              15
                              else:
                                                # Set model to evaluate mode
              16
                                  model.eval()
              17
              18
                              running_loss = 0.0
              19
                              running corrects = 0
              20
              21
                              # Iterate over data.
              22
                              for inputs, labels in dataloaders[phase]:
              23
                                  inputs = inputs.to(device)
              24
                                  labels = labels.to(device)
              25
              26
                                  # zero the parameter gradients
              27
                                  optimizer.zero_grad()
              28
              29
                                  # forward
              30
                                  # track history if only in train
              31
                                  with torch.set_grad_enabled(phase == 'train'):
              32
                                      outputs = model(inputs)
              33
                                       , preds = torch.max(outputs, 1)
              34
                                      loss = criterion(outputs, labels)
              35
              36
                                      # backward + optimize only if in training phase
              37
                                      if phase == 'train':
              38
                                          loss.backward()
              39
                                          optimizer.step()
              40
              41
                                  # statistics
              42
                                  running_loss += loss.item() * inputs.size(0)
              43
                                  running_corrects += torch.sum(preds == labels.data)
              44
                              if phase == 'train':
              45
                                  scheduler.step() # updates the Learning rate
              46
              47
                              epoch_loss = running_loss / dataset_sizes[phase]
              48
                              epoch_acc = running_corrects.double() / dataset_sizes[phase]
              49
              50
                              print('{} Loss: {:.4f} Acc: {:.4f}'.format(
              51
                                  phase, epoch_loss, epoch_acc))
              52
              53
                              # deep copy the model
                              if phase == 'test' and epoch_acc > best_acc:
              54
              55
                                  best_acc = epoch_acc
              56
                                  best model wts = copy.deepcopy(model.state dict())
              57
              58
                          print()
              59
              60
                      time_elapsed = time.time() - since
              61
                      print('Training complete in {:.0f}m {:.0f}s'.format(
              62
                          time_elapsed // 60, time_elapsed % 60))
              63
                      print('Best test Acc: {:4f}'.format(best_acc))
              64
              65
                      # load best model weights
              66
                      model.load_state_dict(best_model_wts)
```

```
In [29]:
               1
                 def train_model(model, criterion, optimizer, scheduler, num_epochs=25):
               2
                      device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
               3
                      model = model.to(device) # Move model to GPU
               4
                      since = time.time()
               5
               6
                      best_model_wts = copy.deepcopy(model.state_dict())
               7
                      best_acc = 0.0
               8
               9
                      for epoch in range(num_epochs):
              10
                          print(f'Epoch {epoch}/{num_epochs - 1}')
              11
                          print('-' * 10)
              12
                          for phase in ['train', 'test']:
              13
                              if phase == 'train':
              14
              15
                                  model.train() # Set model to training mode
              16
                                  model.eval() # Set model to evaluation mode
              17
              18
              19
                              running_loss = 0.0
              20
                              running_corrects = 0
              21
              22
                              for inputs, labels in dataloaders[phase]:
              23
                                  inputs = inputs.to(device) # Move inputs to GPU
                                  labels = labels.to(device) # Move Labels to GPU
              24
              25
              26
                                  optimizer.zero_grad()
              27
              28
                                  with torch.set_grad_enabled(phase == 'train'):
              29
                                      outputs = model(inputs) # Forward pass
              30
                                      _, preds = torch.max(outputs, 1)
              31
                                      loss = criterion(outputs, labels)
              32
                                      if phase == 'train':
              33
              34
                                          loss.backward() # Backpropagation
              35
                                          optimizer.step() # Optimize weights
              36
              37
                                  running_loss += loss.item() * inputs.size(0)
              38
                                  running_corrects += torch.sum(preds == labels.data)
              39
              40
                              if phase == 'train':
              41
                                  scheduler.step()
              42
              43
                              epoch_loss = running_loss / dataset_sizes[phase]
              44
                              epoch_acc = running_corrects.double() / dataset_sizes[phase]
              45
                              print(f'{phase} Loss: {epoch_loss:.4f} Acc: {epoch_acc:.4f}')
              46
              47
              48
                              if phase == 'test' and epoch_acc > best_acc:
              49
                                  best acc = epoch acc
              50
                                  best_model_wts = copy.deepcopy(model.state_dict())
              51
              52
                          print()
              53
              54
                      time_elapsed = time.time() - since
              55
                      print(f'Training complete in {time_elapsed // 60:.0f}m {time_elapsed % 60:.0f}s')
              56
                      print(f'Best test Acc: {best_acc:.4f}')
              57
              58
                      model.load_state_dict(best_model_wts)
              59
                      return model
              60
```

# **Problem 4 (Optional)**

Explore choosing different learning rates or optimizers to see how things go.

#### **Train**

```
In [31]:
          H
              1 # Now we train
               2 | model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler,num_epochs=2
             Epoch 0/24
             ------
             train Loss: 0.4797 Acc: 0.7751
             test Loss: 0.0168 Acc: 1.0000
             Epoch 1/24
             train Loss: 0.4086 Acc: 0.8835
             test Loss: 0.0281 Acc: 0.9841
             Epoch 2/24
             train Loss: 0.2769 Acc: 0.9157
             test Loss: 0.0614 Acc: 0.9683
             Epoch 3/24
             train Loss: 0.2740 Acc: 0.8996
             test Loss: 0.0466 Acc: 0.9841
In [18]:
              1
```

### **Problem 5**

Report your accuracy on the testing set. You should compute this explicitly by running all your testing examples through the model and checking them against their true labels.

```
In [32]:
              1 def evaluate_model(model, dataloader):
                     model.eval() # Set to evaluation mode
              3
                     correct = 0
              4
                     total = 0
              5
                     with torch.no_grad():
                         for inputs, labels in dataloader:
              7
                             inputs, labels = inputs.to(device), labels.to(device)
              8
                             outputs = model(inputs)
              9
                             _, preds = torch.max(outputs, 1)
             10
                             total += labels.size(0)
             11
                             correct += (preds == labels).sum().item()
             12
             13
                     accuracy = correct / total
             14
                     print(f"Test Accuracy: {accuracy:.4f}")
             15
                     return accuracy
             16
             17 evaluate_model(model_ft, dataloaders['test'])
             18
```

Test Accuracy: 1.0000

Out[32]: 1.0

#### Now let's visualize to see how well we did

### **Problem 6**

Print out images from the test set, reporting the predicted labels and the true labels.

```
In [33]:
              1
                 def visualize_predictions(model, dataloader, num_images=6):
               2
                     model.eval()
              3
                     images_so_far = 0
               4
                     fig = plt.figure()
               5
               6
                     with torch.no_grad():
              7
                         for inputs, labels in dataloader:
              8
                             inputs, labels = inputs.to(device), labels.to(device)
              9
                             outputs = model(inputs)
             10
                             _, preds = torch.max(outputs, 1)
             11
             12
                             for i in range(inputs.size()[0]):
             13
                                  images_so_far += 1
             14
                                 ax = plt.subplot(num_images // 2, 2, images_so_far)
             15
                                 ax.axis('off')
                                  ax.set_title(f'True: {class_names[labels[i]]} \nPred: {class_names[pred
             16
             17
                                 imshow(inputs.cpu().data[i])
             18
             19
                                 if images_so_far == num_images:
             20
                                      return
             21
             22 visualize_predictions(model_ft, dataloaders['test'])
              23
```

True: 'vintage\_typewriter'
Pred: 'vintage typewriter'



True: 'antique\_compass' Pred: 'antique compass'



True: 'antique\_compass' Pred: 'antique compass'



True: 'vintage\_typewriter' Pred: 'vintage typewriter'



True: 'antique\_compass' Pred: 'antique compass'



True: 'antique\_compass' Pred: 'antique compass'



# **Problem 7 (Optional)**

Experiment with some/all of the following:

- Play with different pre-trained models
- Experiment with different pre-processing of the data (e.g., turn data augmentation on or off).
- Experiment with adding more/fewer layers, and/or layers of different size.
- Fine tune by training everything or only the last (new) layer -- this requires figuring out how to only update some of the layers.
- Try to reduce the number of images you use for training. How few can you use and still get good accuracy? Remember that you were all able to learn what a Goblin shark is with only one single example.